




H 3

Computer Arithmetic

ELECTA01 **Computer architecture**

ELECTA01
Hogeschool Rotterdam
Opleiding Elektrotechniek
Minor Embedded Systems
J.W. Peltenburg / J.Z.M. Broeders (brojz@hr.nl)

ARITHMETIC FOR COMPUTERS

Arithmetic for Computers

- Operations on integers
 - Addition and subtraction
 - Multiplication and division
 - Dealing with overflow
- Floating-point real numbers
 - Representation and operations


Addition and Subtraction



- Right to left addition, bit by bit.
- Carries passed to the next digit to the left.
- Behavior as expected, because we humans do the same.
- Subtraction is the same as addition: the appropriate operand is simply negated before being added.

Integer Addition


■ Example: $7 + 6$



■ Overflow if result out of range

- Adding +ve and -ve operands, no overflow
- Adding two +ve operands
 - Overflow if result sign is 1
- Adding two -ve operands
 - Overflow if result sign is 0

Integer Subtraction



- Add negation of second operand
- Example: $7 - 6 = 7 + (-6)$

+7:	0000 0000 ... 0000 0111
-6:	1111 1111 ... 1111 1010
+1:	0000 0000 ... 0000 0001

- Overflow if result out of range
 - Subtracting two +ve or two -ve operands, no overflow
 - Subtracting +ve from -ve operand
 - Overflow if result sign is 0
 - Subtracting -ve from +ve operand
 - Overflow if result sign is 1

Overflows




- Can we have an overflow by adding if both the operands have different signs?


Operation	Operand A	Operand B	Result indicating overflow
$A + B$	≥ 0	≥ 0	< 0
$A + B$	< 0	< 0	≥ 0
$A - B$	≥ 0	< 0	< 0
$A - B$	< 0	≥ 0	≥ 0

Multiplication



- Start with long-multiplication approach



Length of product is the sum of operand lengths



Multiplication Hardware



Corrected version of Figure 3.4 page 194.

Exercise


- Multiply 10 by 9
- Use 4 bits multiplicand and 4 bits multiplier (8 bits product).

Exercise solution

Iteration	Step	Multiplier	Multiplicand	Product
0	Initial values	1001	0000 1010	0000 0000
1	multiplier(0)=1; add multiplicand sll multiplicand slr multiplier	1001	0000 1010	0000 1010
		1001	0001 0100	0000 1010
		0100	0001 0100	0000 1010
2	multiplier(0)=0; no add multiplicand sll multiplicand slr multiplier	0100	0001 0100	0000 1010
		0100	0010 1000	0000 1010
		0010	0010 1000	0000 1010
3	multiplier(0)=0; no add multiplicand sll multiplicand slr multiplier	0010	0010 1000	0000 1010
		0010	0101 0000	0000 1010
		0001	0101 0000	0000 1010
4	multiplier(0)=1; add multiplicand sll multiplicand slr multiplier	0001	0101 0000	0101 1010
		0001	1010 0000	0101 1010
		0000	1010 0000	0101 1010

Optimized Multiplier


- Perform steps in parallel: add/shift



- One cycle per partial-product addition
 - That's ok, if frequency of multiplications is low

Faster Multiplier

- Uses multiple adders
 - Cost/performance tradeoff



- Can be pipelined
 - Several multiplication performed in parallel

LEGv8 Multiplication

- Three multiply instructions:
 - MUL: multiply
 - Gives the lower 64 bits of the product
 - SMULH: signed multiply high
 - Gives the upper 64 bits of the product, assuming the operands are signed
 - UMULH: unsigned multiply high
 - Gives the upper 64 bits of the product, assuming the operands are unsigned

Floating Point

- Representation for non-integral numbers
 - Including very small and very large numbers
- Like scientific notation
 - -2.34×10^{56} ← normalized
 - $+0.002 \times 10^{-4}$ ← not normalized
 - $+987.02 \times 10^9$ ← not normalized
- In binary
 - $\pm 1.xxxxxxx_2 \times 2^{yyyy}$
- Types `float` and `double` in C

Floating Point Standard

- Defined by IEEE Std 754-1985
 - latest version IEEE 754-2008
- Developed in response to divergence of representations
 - Portability issues for scientific code
- Now almost universally adopted
- Several representations, most importantly:
 - Single precision (32-bit)
 - Double precision (64-bit)

IEEE Floating-Point Format

single: 8 bits

double: 11 bits

single: 23 bits

double: 52 bits

S	Exponent	Fraction
---	----------	----------

$$x = (-1)^S \times (1.\text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0 $\Rightarrow$ non-negative, 1 $\Rightarrow$ negative)
- Normalize significand: $1.0 \leq |\text{significand}| < 2.0$
 - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
 - Significand is Fraction with the “1.” restored
- Exponent: excess representation: actual exponent + Bias
 - Ensures exponent is unsigned
 - Single: Bias = 127; Double: Bias = 1023

Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
 - Exponent: 00000001
 $\Rightarrow$ actual exponent = $1 - 127 = -126$
 - Fraction: 000...00 $\Rightarrow$ significand = 1.0
 - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$
- Largest value
 - exponent: 11111110
 $\Rightarrow$ actual exponent = $254 - 127 = +127$
 - Fraction: 111...11 $\Rightarrow$ significand ≈ 2.0
 - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

Double-Precision Range

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
 - Exponent: 000000000001
 $\Rightarrow$ actual exponent = $1 - 1023 = -1022$
 - Fraction: 000...00 $\Rightarrow$ significand = 1.0
 - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- Largest value
 - Exponent: 111111111110
 $\Rightarrow$ actual exponent = $2046 - 1023 = +1023$
 - Fraction: 111...11 $\Rightarrow$ significand ≈ 2.0
 - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

Floating-Point Precision

- Relative precision
 - all fraction bits are significant
 - Implicit 1 is significant
 - **Single**: approx 2^{-24}
 - Equivalent to $24 \times \log_{10} 2 \approx 24 \times 0.3 \approx 7$ decimal digits of precision
 - **Double**: approx 2^{-53}
 - Equivalent to $53 \times \log_{10} 2 \approx 53 \times 0.3 \approx 16$ decimal digits of precision

Floating-Point Example

- Represent -0.75 in IEEE 754 single and double precision
 - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
 - $S = 1$
 - Fraction = $1000\dots00_2$
 - Exponent = $-1 + \text{Bias}$
 - Single: $-1 + 127 = 126 = 01111110_2$
 - Double: $-1 + 1023 = 1022 = 01111111110_2$
- Single: $10111111101000\dots00$ $0x\text{BF}400000$
- Double: $101111111111101000\dots00$

Floating-Point Example

- What number is represented by the single precision float (IEEE 754) 0xC0A00000

11000000101000...00

- $S = 1$

- Fraction = $01000...00_2$

- Exponent = $10000001_2 = 129$

- $$\begin{aligned} x &= (-1)^1 \times (1.01_2) \times 2^{(129 - 127)} \\ &= (-1) \times 1.25 \times 2^2 \\ &= -5.0 \end{aligned}$$

Denormal Numbers

- Exponent = 000...0 $\Rightarrow$ hidden bit is 0

$$x = (-1)^S \times (0.\text{Fraction}) \times 2^{-\text{Bias}}$$

- Smaller than normal numbers
 - allow for gradual underflow, with diminishing precision

- Denormal with fraction = 000...0

$$x = (-1)^S \times (0.0) \times 2^{-\text{Bias}} = \pm 0.0$$

Two representations
of 0.0!



Infinites and NaNs

- Exponent = 111...1, Fraction = 000...0
 - $\pm$ Infinity
 - Can be used in subsequent calculations, avoiding need for overflow check
- Exponent = 111...1, Fraction $\neq$ 000...0
 - Not-a-Number (NaN)
 - Indicates illegal or undefined result
 - e.g., 0.0 / 0.0
 - Can be used in subsequent calculations

FP Instructions in LEGv8

- Separate FP registers
 - 32 x 32-bits single-precision: S0, ..., S31
 - 32 x 64-bits double-precision: D0, ..., D31
 - S_n stored in the lower 32 bits of D_n
- FP instructions operate only on FP registers
 - Programs generally don't do integer ops on FP data, or vice versa
 - More registers with minimal code-size impact
- FP load and store instructions
 - LDURS, LDURD
 - STURS, STURD

FP Instructions in LEGv8

- Single-precision arithmetic
 - FADDS, FSUBS, FMULS, FDIVS
 - e.g., FADDS S2, S4, S6
- Double-precision arithmetic
 - FADDD, FSUBD, FMULD, FDIVD
 - e.g., FADDD D2, D4, D6
- Single- and double-precision comparison
 - FCMPS, FCMPD
 - Sets or clears FP condition-code bits
- Branch on FP condition code true or false
 - B.cond

FP Example: °F to °C

- C code:

```
float f2c (float fahr) {  
    return ((5.0/9.0)*(fahr - 32.0));  
}
```

- fahr in S0, result in S0, literals in global memory space with offset from X27 (const5, const9, const32).

- Compiled LEGv8 code:

f2c:

```
LDURS S16, [X27, const5]  
LDURS S17, [X27, const9]  
FDIVS S18, S16, S17  
LDURS S19, [X27, const32]  
FSUBS S20, S0, S19  
FMULS S0, S18, S20  
BR     LR
```

ARM Call Convention:
D0-D7: Arguments / Results
D8-D15: Saved by callee
D16-D31: Temporaries

FP Example: Matrix Multiplication

- $X = X + Y \times Z$
 - All 32×32 matrices, 64-bit double-precision elements

- C code:

```
void mm(double x[][],
        double y[][], double z[][]) {
    int i, j, k;
    for (i = 0; i != 32; i = i + 1)
        for (j = 0; j != 32; j = j + 1)
            for (k = 0; k != 32; k = k + 1)
                x[i][j] = x[i][j]
                    + y[i][k] * z[k][j];
}
```

- Addresses of x, y, z in X0, X1, X2, and
i, j, k in X9, X10, X11

FP Example: Matrix Multiplication

■ LEGv8 code:

```
mm:  ADDI  X12, XZR, #32
      MOV   X9, XZR
L1:   MOV   X10, XZR
L2:   MOV   X11, XZR
      LSL   X13, X9, #5
      ADD   X14, X13, X10
      LSL   X15, X14, #3
      ADD   X16, X0, X15
      LDURD D8, [X16, #0]
L3:   LSL   X13, X9, #5
      ADD   X14, X13, X11
      LSL   X15, X14, #3
      ADD   X17, X1, X15
      LDURD D9, [X17, #0]
```

```
LSL   X13, X11, #5
ADD    X14, X13, X11
LSL    X15, X14, #3
ADD    X17, X2, X15
LDURD  D10, [X17, #0]
FMULD  D11, D9, D10
FADDD  D8, D8, D11
ADDI   X11, X11, #1
CMP    X11, X12
B.LT   L3
STURD  D8, [X16, #0]
ADDI   X10, X10, #1
CMP    X10, X12
B.LT   L2
ADDI   X9, X9, #1
CMP    X9, X12
B.LT   L1
BR     LR
```

Concluding Remarks

- Bits have no inherent meaning
 - Interpretation depends on the instructions applied
- Computer representations of numbers
 - Finite range and precision
 - Need to account for this in programs

Concluding Remarks

- ISAs support arithmetic
 - Signed and unsigned integers
 - Floating-point approximation to reals
- Bounded range and precision
 - Operations can overflow and underflow